GLOSS-MT: Controllable Low-Resource Machine Translation with an Editable Dictionary

Mich-Seth Owusu Stephen E. Moore Akwasi Asare Lawrence Adu Gyamfi
Paul Azunre Joel Budu Jonathan Asiamah

Correspondence: michseth@ghananlp.org

Abstract

Adapting a pretrained machine translation (MT) model to new terminology usually requires new parallel data and another round of fine-tuning, which is costly for low-resource languages. We ask whether controllability can be added to MT, and the fine-tuning burden reduced, without degrading the pretrained model. GLOSS-MT looks up the source's noun phrases in an offline English–Twi dictionary, replaces every phrase found by its Twi translation in place, and fine-tunes a translation model once to translate these mixed inputs. We apply the same data and pipeline to NLLB-200 (600M, 1.3B, 3.3B), OPUS-MT en-tw and TranslateGemma 4B. On a document-level English→Asante Twi benchmark (84 documents, four domains), every fine-tune scores above its pretrained base in chrF2 (+2.85 to +5.48), and the three NLLB fine-tunes are the best open systems (50.12–50.80). A direct controllability test on human-translated Bible passages shows that the three fine-tunes tested write 79–97% of the dictionary terms, and that replacing every dictionary entry by a different term changes the output accordingly (70–97% of the new terms appear) with no retraining. We release the models, data and code.

1 Introduction

Twi (Akan) has tens of millions of speakers but little parallel text. Pretrained multilingual models such as NLLB-200 [10] cover it, and fine-tuning them on a few thousand in-domain pairs is known to help [1]. But each new domain or terminology convention (a ministry's budget terms, legal vocabulary, a dialectal preference) then needs new parallel data and new training. Bilingual wordlists and glossaries, in contrast, are the resource that language communities and field linguists most often produce.

We study the question: *can we add controllability to machine translation, and reduce the burden of fine-tuning, without degrading the performance of pretrained MT models?* It has three parts: (i) **control**: can an operator change which term a phrase receives by editing an external dictionary instead of retraining? (ii) **a light fine-tune**: can one short fine-tune teach a model to use dictionary terms, so that later changes need only dictionary entries? (iii) **no degradation**: does the model translate at least as well as the pretrained model it starts from?

GLOSS-MT answers with a simple pipeline: dictionary terms are pretranslated *in place* in the source, and the model is fine-tuned once to keep them. Our contributions are: (1) the pipeline and a training set built by aligning English noun phrases to their spans in gold Twi; (2) fine-tunes of five pretrained models of three families, each compared with its own base on an independent document-level benchmark; (3) a direct controllability test with a counterfactual ("swap") dictionary, on human-translated references; and (4) released models, data and code.[1]

2 Related Work

Terminology can be imposed at decoding time with lexically constrained beam search [6, 14], which guarantees the terms but slows decoding and handles inflection poorly. Alternatively, the model is trained to use terms annotated in its input: Dinu et al. [4] append target terms inline next to the source words, Song et al. [16] replace source phrases by their prespecified translations (code-switching), and Bergmanis and Pinnis [3] annotate target lemmas so the model can inflect them. Ghazvininejad et al. [5] prompt LLMs with dictionary entries for rare words. GLOSS-MT follows the in-place replacement idea of Song et al. [16], applies it with a large community-style dictionary rather than a few constraints per sentence, and asks a different question: whether a *single* fine-tune of an existing pretrained MT model makes the dictionary an operator's control surface without costing translation quality. For Twi, prior work built English–Twi parallel data [2] and fine-tuned pretrained models on news translation [1].

3 Method

Dictionary. An offline English→Twi dictionary of 613,259 phrases with exactly one translation each. Entries come from machine-translated phrase lists (re-translated with Gemini where an entry had several candidates) and from the phrase–span alignments of the training data below.

[1] Models: https://huggingface.co/ghanaopenai; code: https://github.com/GhanaOpenAI/gloss-mt; benchmark: https://github.com/GhanaNLP/nsanku-mt-benchmark.

Pipeline. A document is split into lines, and consecutive lines are packed into chunks of at most 192 source tokens (128 for OPUS-MT), cutting only at line boundaries so that Markdown structure is preserved. The chunk's noun phrases are extracted with spaCy [7] and looked up; each phrase found is replaced in the English, longest first, by its Twi translation in angle brackets, and phrases not found stay English. For example, "*Our people have courage, but courage without foundation is foolishness.*" becomes "*Our <nnipa> have <akokoduro>, but <akokoduro> without <fapem> is <nkwasiasɛm>.*" The model writes the terms in brackets too, and the pipeline strips them. Table-of-contents dot leaders are kept out of the model and restored.

Training data. From 15,134 English–Twi training documents we build 31,909 line-aligned chunks. Gemini aligns each English noun phrase of a chunk to its span in the reference Twi; a span is kept only if it occurs verbatim on the phrase's line. Each aligned pair is bracketed on both sides: the input carries the Twi span in place of the phrase, and the target carries the same span in brackets. The model thus learns to copy the bracketed terms and translate the sentence around them. The data is released as `gloss-mt-parallel-en-twi`.

Models. We fine-tune NLLB-200 3.3B with LoRA [8] (r=32 on all attention and feed-forward projections, 2% of the parameters, merged after training); NLLB-200 1.3B and distilled 600M with full fine-tuning (Adafactor [15], lr 10^{-4}); OPUS-MT en-tw [17] ($\sim$75M parameters, full fine-tuning, AdamW, lr 10^{-4}); and TranslateGemma 4B,[2] a translation-specialised Gemma 3 model, with LoRA (r=32), using its own prompt with target code `ak` since it has no Twi code. A literal new-line token is added to the NLLB tokenizer. Every 1,000 steps the model translates 200 fixed validation chunks, and the checkpoint with the best mean of BLEU, chrF2 and $100\,e^{-\text{loss}}$ is kept, with early stopping. NLLB and OPUS-MT decode with beam 4, TranslateGemma greedily.

4 Evaluation

Benchmark. The nsanku MT benchmark scores document-level chrF2 [12] and BLEU [11] (sacreBLEU [13]) over 84 English→Asante Twi documents in four datasets: three Citizens' Budget documents (Finance, 2021–2023), the Universal Declaration of Human Rights (UDHR), five Watchtower study editions (2026), and 75 Akan folktales (Rattray, 1930). Because 75 documents are folktales, the language score is the mean of the four per-dataset scores. Each fine-tune is compared with its pretrained base, translated with the same chunking; commercial systems are references.

[2] `https://huggingface.co/google/translategemma-4b-it`

Controllability. Each test set is translated three times by the same model, changing only the dictionary: *dict* (as deployed), *none* (no dictionary) and *swap*, a counterfactual dictionary in which every phrase found in the set receives a different real Twi term from the set's other terms, chosen so that it occurs neither in the source nor in the dict or none outputs of that passage; a swapped term can therefore only reach the output through the dictionary. A term counts as used when its words appear in order as whole words of the output. We report **adherence** (share of inserted terms in the output), **natural rate** (the same terms without the dictionary), **swap adherence** (share of swapped-in terms in the swap output) and **leak** (share of original terms surviving the swap). The main test set is 300 passages of three consecutive Bible verses, with the Good News Translation as English and the human Asante Twi *Nkwa Asɛm* as reference (2,766 inserted terms); no passage shares an 8-word span with the training data. Both versions are free, meaning-based translations, so absolute scores are low and the with/without comparison is what matters. We also run the test on 300 validation chunks (5,387 terms) whose references are not all human-curated.

5 Results

No degradation in chrF2. Table 1 shows that every fine-tune scores above its pretrained base in chrF2, by +2.85 to +5.48. The gain concentrates in contemporary prose: Watchtower rises by 5.7–7.3 points for the NLLB models and Finance by 3.7–5.7. The only domain below its base is the UDHR for NLLB 3.3B (-0.39). The 1930 folktales, written in pre-standard orthography, gain least. The three NLLB fine-tunes are the best open systems, above Khaya AI and behind only Gemini 3.8 Flash and Google Translate. OPUS-MT gains most (+5.48) and, at 75M parameters, scores above the base NLLB 3.3B (48.00 vs 47.74).

Model size matters little once fine-tuned. The fine-tuned NLLB models span only 50.12–50.80 from 600M to 3.3B, and the 600M fine-tune scores 2.38 above the base 3.3B model, so the cheapest model is a practical deployment choice.

BLEU. The NLLB fine-tunes are slightly below their bases in BLEU (-0.51 to -1.04), mostly on the UDHR and Watchtower, while chrF2 rises on both; OPUS-MT gains BLEU (+1.20). The fine-tuned output matches more of the reference's characters but fewer of its exact word sequences, plausibly because dictionary terms differ from the reference's word choices. We report this as a limit on the no-degradation claim.

Validation scores do not predict the benchmark. TranslateGemma scored above NLLB 3.3B on validation (chrF2 64.96 vs 63.66) but below it on the benchmark, and the 600M model above the 1.3B on validation (63.99 vs 62.14) but slightly below it on the benchmark. Checkpoint selection on

	System	chrF2	BLEU	Finance	UDHR	Watchtower	Folktales
References	Gemini 3.8 Flash	54.64	17.25	52.50	56.44	69.28	40.33
	Google Translate	51.26	15.67	48.06	52.30	67.15	37.54
	Khaya AI	49.91	12.79	47.74	49.92	65.78	36.22
NLLB-200 3.3B	base	47.74	**14.15**	45.25	**52.05**	56.89	36.76
	GLOSS-MT	**50.80** (+3.06)	13.11	**50.23**	51.66	**64.17**	**37.16**
NLLB-200 1.3B	base	47.42	**13.68**	45.72	50.47	57.46	36.05
	GLOSS-MT	**50.27** (+2.85)	12.82	**49.41**	**51.51**	**63.66**	**36.51**
NLLB-200 600M	base	46.71	**12.74**	44.10	50.17	57.06	35.50
	GLOSS-MT	**50.12** (+3.41)	12.23	**49.81**	**51.25**	**62.80**	**36.61**
OPUS-MT en-tw	base	42.52	9.71	37.78	47.52	54.23	30.53
	GLOSS-MT	**48.00** (+5.48)	**10.91**	**46.13**	**50.09**	**60.62**	**35.15**
TranslateGemma 4B	GLOSS-MT	49.81	12.24	50.03	51.19	62.54	35.48

Table 1: English→Asante Twi on the nsanku benchmark (84 documents; chrF2 per dataset; the mean is over the four datasets). Bold: better of base and GLOSS-MT. Zero-shot, TranslateGemma does not produce Twi, so it has no base row. MADLAD-400 3B [9], the remaining open baseline, scores 37.14.

Model	Adh.	Nat.	Swap	Leak	chrF2 d/n
Asante Twi Bible (human references)					
GLOSS-MT TranslateGemma	**96.6**	26.2	**97.3**	1.9	37.2 / 32.1
GLOSS-MT NLLB 3.3B	94.4	38.4	95.2	2.2	41.2 / 40.9
GLOSS-MT OPUS-MT	79.2	32.5	69.7	1.7	36.4 / 33.6
NLLB 3.3B base	66.0	33.9	72.0	2.8	38.9 / 39.6
OPUS-MT base	1.9	30.3	0.4	1.4	19.1 / 30.1
Validation chunks					
GLOSS-MT TranslateGemma	**97.8**	36.1	**95.9**	6.1	66.0 / 67.7
GLOSS-MT NLLB 3.3B	88.4	37.8	75.8	6.6	65.4 / 71.0
GLOSS-MT OPUS-MT	76.2	34.7	61.0	5.6	59.2 / 64.3
NLLB 3.3B base	55.8	26.9	49.0	3.6	41.3 / 45.5
OPUS-MT base	1.8	13.9	0.1	1.3	15.0 / 26.1

Table 2: Controllability (%). Adh.: inserted terms in the output; Nat.: the same terms without the dictionary; Swap: swapped-in terms in the output; Leak: original terms surviving the swap; chrF2 with / without the dictionary.

in-domain validation data is a weak proxy for document-level quality.

The dictionary controls the fine-tuned models. We test three fine-tunes, one per model family, and two of their bases. On the Bible (Table 2), GLOSS-MT TranslateGemma and NLLB 3.3B write 95–97% of the inserted terms, against 26–38% of the same terms without the dictionary. After every entry is replaced, they write 95–97% of the new terms and about 2% of the original ones survive: editing the dictionary changes the translation as intended, without retraining. TranslateGemma follows multi-word terms as reliably as single words (98.1% vs 97.7% on validation). OPUS-MT, the smallest model, follows 70% of swapped terms; the swapped terms are often wrong in context, and a small model drops them more often.

Fine-tuning creates the control. Base OPUS-MT ignores the bracketed terms (1.9%) and degrades when they are present (chrF2 30.1→19.1). Base NLLB 3.3B copies some Twi from its input (66.0%), and fine-tuning raises adherence to 94.4%.

No quality cost on human references. On the Bible, translating with the dictionary scores at least as well in chrF2 as without it: TranslateGemma +5.0, OPUS-MT +2.8, NLLB 3.3B +0.3 (BLEU 10.75 vs 10.91). This holds although only 35.9% of the inserted terms occur in the human reference. On the validation chunks the dictionary instead *lowers* chrF2 for every fine-tune (-1.7 to -5.6); the fine-tuned models reproduce those references' style more closely without the dictionary, which suggests the drop reflects the references rather than a loss of quality, and is why we rely on the human-referenced set for this comparison.

6 Limitations

Control is learned, not enforced: adherence is high but not complete, and lower for the smallest model. Term matching is by exact whole words, so a term that Twi morphology changes (a prefix or plural) counts as missed and adherence is a lower bound; target-lemma annotation [3] could address inflection. The dictionary is largely machine-generated, and its entries often differ from the human references' choices. The fine-tuned NLLB models lose some BLEU. The human-referenced controllability set is a single domain (the Bible), translated freely on both sides, and we have not yet measured the quality gain from adding curated terminology for a new domain. All evaluation is English→Asante Twi.

7 Conclusion

A single fine-tune on dictionary-pretranslated inputs makes an external dictionary an effective control surface for pretrained MT models: the three fine-tunes we tested follow 79–97% of the dictionary terms and follow edited entries without

retraining, and all five fine-tunes, from three model families, score above their pretrained bases in chrF2 on an independent document-level benchmark, with a small BLEU cost for the NLLB models. Because the gain is similar from 600M to 3.3B parameters, communities can obtain controllable translation from a small model and maintain it by editing wordlists rather than collecting parallel data. Models, data and code are released.

References

[1] David Ifeoluwa Adelani et al. A few thousand translations go a long way! leveraging pre-trained models for African news translation. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 3053–3070, 2022.

[2] Paul Azunre et al. English-Twi parallel corpus for machine translation. *arXiv preprint arXiv:2103.15625*, 2021.

[3] Toms Bergmanis and Mārcis Pinnis. Facilitating terminology translation with target lemma annotations. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, pages 3105–3111, 2021.

[4] Georgiana Dinu, Prashant Mathur, Marcello Federico, and Yaser Al-Onaizan. Training neural machine translation to apply terminology constraints. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3063–3068, 2019.

[5] Marjan Ghazvininejad, Hila Gonen, and Luke Zettlemoyer. Dictionary-based phrase-level prompting of large language models for machine translation. *arXiv preprint arXiv:2302.07856*, 2023.

[6] Chris Hokamp and Qun Liu. Lexically constrained decoding for sequence generation using grid beam search. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics*, pages 1535–1546, 2017.

[7] Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spaCy: Industrial-strength natural language processing in Python, 2020.

[8] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.

[9] Sneha Kudugunta, Isaac Caswell, Biao Zhang, Xavier Garcia, Christopher A. Choquette-Choo, Katherine Lee, Derrick Xin, Aditya Kusupati, Romi Stella, Ankur Bapna, and Orhan Firat. MADLAD-400: A multilingual and document-level large audited dataset. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2023.

[10] NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, et al. No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*, 2022.

[11] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, 2002.

[12] Maja Popović. chrF: character n-gram F-score for automatic MT evaluation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, 2015.

[13] Matt Post. A call for clarity in reporting BLEU scores. In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186–191, 2018.

[14] Matt Post and David Vilar. Fast lexically constrained decoding with dynamic beam allocation for neural machine translation. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 1314–1324, 2018.

[15] Noam Shazeer and Mitchell Stern. Adafactor: Adaptive learning rates with sublinear memory cost. In *Proceedings of the 35th International Conference on Machine Learning*, pages 4596–4604, 2018.

[16] Kai Song, Yue Zhang, Heng Yu, Weihua Luo, Kun Wang, and Min Zhang. Code-switching for enhancing NMT with pre-specified translation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 449–459, 2019.

[17] Jörg Tiedemann and Santhosh Thottingal. OPUS-MT – building open translation services for the world. In *Proceedings of the 22nd Annual Conference of the European Association for Machine Translation*, pages 479–480, 2020.